TARGET-TO-COMMITTEE DISAGREEMENT FOR MRI SLICE REVIEW: SELECTION, TRANSFER, AND FAILURE BOUNDARIES

Anonymous Authors

Anonymous submission

ABSTRACT

Cohort-average fidelity cannot identify which MRI interpolation should be reviewed when the missing target is unavailable. We evaluate ground-truth-free ranking from deterministic outputs: a target interpolator is excluded from a heterogeneous committee, and its median target-to-member difference is aggregated over generated foreground voxels. Ground truth is isolated from scoring. We register 132 score–mask–pooling candidates and 40 score–committee combinations on IXI $\times 4$ and fastMRI-brain $\times 2/\times 4/\times 6$, freeze four finalists, and validate once on 22 eight-fold out-of-fold NingXia examinations at all factors. A development-selected committee reduces mean nAURC from 0.089 to 0.040 but degrades it from 0.070 to 0.116 on frozen validation. The pre-specified score remains strongest (Spearman 0.901, nAURC 0.070, RRR50 21.8%). This transfer failure, paired bootstrap analysis, and common-mode decomposition show that committee selection can overfit cohort difficulty. Disagreement supports selective review, not calibrated uncertainty or clinical certification.

***Index Terms*—** MRI interpolation, reliability ranking, selective review, model disagreement, quality control

1. INTRODUCTION

Through-plane interpolation densifies sparse MRI stacks for visualization, registration, and volumetric processing. Modern methods model inter-slice context or adapt restoration transformers [1, 2, 3], yet they are usually compared by cohort-average PSNR, SSIM, or MAE. Such metrics summarize a test set but do not identify an unreliable examination once the missing slice is unobservable.

Model-internal uncertainty methods use stochastic inference, repeated training, heteroscedastic objectives, TTA, posterior reconstruction, or learned error prediction [4, 5, 6, 7, 8, 9, 10, 11, 12]; no-reference MRI quality assessment instead targets perceptual quality or acquisition artifacts [13]. Our setting assumes only deterministic outputs. Selective prediction supplies the endpoint [14, 15, 16, 17]. We contribute (1) a self-excluded target-to-committee ranking protocol and common-mode error boundary, (2) a registered candidate funnel with strict development/validation separation, and (3) a multi-factor transfer study showing that development-selected committees need not improve independent-cohort review.

2. METHOD

2.1. Deployment contract and target isolation

Let x denote observed sparse-slice context, y the unobserved dense target, and $\hat{y}_m = f_m(x)$ the output of target interpolator m. Available output identities form $\mathcal{B}$, and the target-specific committee is

$$\mathcal{C}_m = \{\hat{y}_b : b \in \mathcal{B},\ b \neq m\}. \tag{1}$$

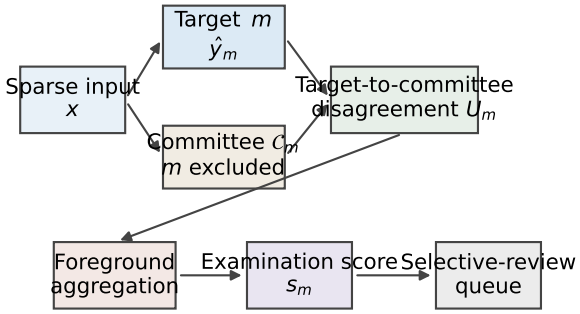


Fig. 1. Target-specific scoring and offline evaluation. The target is excluded from its heterogeneous committee; ground truth is revealed only after scoring.

The committee may combine neural point reconstructions with deterministic analytical interpolation, but never contains y or $\hat{y}_m$ itself. Before scoring, outputs must share case identity, affine, orientation, voxel grid, shape, generated-slice pattern, and source-intensity representation. Missing, duplicate, non-finite, or geometrically mismatched files invalidate the profile rather than being silently resampled. Disagreement is meaningful only under this deployment contract.

2.2. Target-specific disagreement family

For generated voxel i, the default map is

$$U_m(i) = \underset{b \in \mathcal{C}_m}{\text{median}} \, |\hat{y}_m(i) - \hat{y}_b(i)|. \tag{2}$$

This is a median target-to-committee difference, not statistical median absolute deviation and not posterior variance. Self-exclusion is essential because including $\hat{y}_m$ creates a structural zero. To separate the protocol from the elementary aggregator, define $d_{mb}(i) = |\hat{y}_m(i) - \hat{y}_b(i)|$ and

$$U_m^g(i) = g(\{d_{mb}(i) : b \in \mathcal{C}_m\}), \tag{3}$$

where g may be the median, mean, maximum, or trimmed mean. Comparators also include distance to the committee mean/median image, target-excluded committee spread, mean/median pairwise disagreement, residuals to linear/cubic interpolation, and gradient- or Laplacian-domain variants. These tests distinguish robust target conditioning from generic set dispersion, analytical departure, and local structural inconsistency. Changing the named target changes every d_{mb} and may change the review order; a single all-output spread map cannot attribute risk to one interpolator.

2.3. Generated support, masking, and intensity

For factor r, let $\mathcal{Z}_{\text{obs}}^{(r)}$ be the recorded anchor indices and $\mathcal{Z}_{\text{gen}}^{(r)} = \mathcal{Z} \setminus \mathcal{Z}_{\text{obs}}^{(r)}$. The audited runs begin at anchor $a = 0$, hence $\mathcal{Z}_{\text{obs}}^{(r)} = \{z : (z - a) \bmod r = 0\}$; define $\mathcal{G}_r = \{i : i_z \in \mathcal{Z}_{\text{gen}}^{(r)}\}$. Observed anchors are never scored. The default deployment mask and examination score are

$$M_m^{(r)}(i) = \mathbf{1}\left[|\hat{y}_m(i)| > 0.01 \max_{j \in \mathcal{G}_r} |\hat{y}_m(j)|\right], \tag{4}$$

$$s_m^{(r)} = \frac{\sum_{i \in \mathcal{G}_r} M_m^{(r)}(i) U_m(i)}{\sum_{i \in \mathcal{G}_r} M_m^{(r)}(i)}. \tag{5}$$

The formulation is independent of factor; r only determines $\mathcal{G}_r$. Each factor forms a separate review queue because raw score scales need not be comparable. Foreground aggregation prevents background dominance, although a target mask may omit suppressed tissue. Ground-truth masks are rejected by the deployment path. Predictions share the source-volume intensity representation; independent output normalization is only a sensitivity because it can erase cross-model amplitude information.

2.4. Identifiability and common-mode invariance

Write $\hat{y}_b = y + c + \epsilon_b$, where c is shared error and ϵ_b is member-specific. Then

$$\hat{y}_m - \hat{y}_b = \epsilon_m - \epsilon_b. \tag{6}$$

For any perturbation c' applied identically to the target and every committee member on the evaluated support,

$$U_m(\{\hat{y}_b + c'\}) = U_m(\{\hat{y}_b\}), \tag{7}$$

because each target-to-member difference is unchanged. This exact common-mode invariance is both the method's rationale and its blind spot: target-specific separation can be ranked, but shared deletion, blur, bias, or preprocessing failure is unidentifiable from disagreement alone. Committee diversity is therefore a necessary operating condition, not a guarantee that consensus is correct.

2.5. Selective review and implementation separation

Ground truth enters only offline. Let E_j be foreground MAE for examination j. Sorting by increasing s_m yields the retained set A_q at coverage q and risk

$$R(q) = |A_q|^{-1} \sum_{j \in A_q} E_j. \tag{8}$$

We report foreground pixel Pearson correlation, examination-level Spearman ρ, absolute retained MAE, $\text{RRR}(q) = [R_{\text{random}} - R(q)]/R_{\text{random}}$, and

$$\text{nAURC} = \frac{\text{AURC} - \text{AURC}_{\text{oracle}}}{\text{AURC}_{\text{random}} - \text{AURC}_{\text{oracle}}}, \tag{9}$$

where zero is oracle and one expected random review. Coverage is always the retained fraction; $1 - q$ is referred. Only ordering is interpreted—the magnitude of s_m is not an error probability. Computing Eq. (2) costs $O(|\mathcal{G}||\mathcal{C}_m|)$ and can stream members. The scoring API receives predictions and metadata only; ground truth, oracle ranking, and error-derived labels reside in a separate evaluator, preventing accidental leakage.

Table 1. Cohorts and stored neural outputs. Analytical references are available for every profile.

Cohort	n	Neural target outputs	Evaluation design
IXI	115	I^{3}Net, HAT, SwinIR, ACSI	fixed held-out outputs
fastMRI-brain	24	I^{3}Net, HAT, SwinIR	fixed stored outputs
NingXia	22	I^{3}Net, HAT, SwinIR	eight-fold out-of-fold

Table 2. All ten pre-specified profiles. Lower nAURC is better; coverage is retained fraction.

Profile	Pixel r_z	Case ρ [95% CI]	nAURC [95% CI]	RRR_{50}
IXI/ACSI	0.435	0.954 [0.925, 0.972]	0.035 [0.018, 0.060]	50.1%
IXI/HAT	0.512	0.966 [0.940, 0.980]	0.024 [0.012, 0.042]	49.6%
IXI/I^{3}Net	0.351	0.951 [0.922, 0.969]	0.044 [0.024, 0.073]	52.1%
IXI/SwinIR	0.520	0.965 [0.937, 0.980]	0.025 [0.012, 0.044]	49.5%
fastMRI/HAT	0.328	0.916 [0.758, 0.962]	0.067 [0.017, 0.176]	19.7%
fastMRI/I^{3}Net	0.291	0.890 [0.713, 0.966]	0.092 [0.029, 0.194]	19.3%
fastMRI/SwinIR	0.363	0.931 [0.781, 0.978]	0.061 [0.012, 0.150]	20.5%
NingXia/HAT	0.277	0.929 [0.790, 0.975]	0.076 [0.016, 0.228]	22.2%
NingXia/I^{3}Net	0.356	0.914 [0.769, 0.970]	0.063 [0.016, 0.186]	23.1%
NingXia/SwinIR	0.279	0.936 [0.807, 0.983]	0.036 [0.007, 0.120]	21.4%

3. EXPERIMENTS

3.1. Cohorts, outputs, and frozen protocol

The factor-four study uses IXI (115 examinations) [18], fastMRI-brain (24) [19], and an internal NingXia cohort (22) pooled from eight held-out folds. IXI includes I^{3}Net, HAT, SwinIR, and ACSI; the other cohorts include the first three. Linear and cubic interpolation provide analytical references [20, 21, 22], yielding 161 examinations and ten fixed profiles. The pre-declared cross-factor extension uses the identical fastMRI-brain and NingXia identities at $r \in \{2, 4, 6\}$ and freezes the common model set to I^{3}Net, HAT, SwinIR, linear, and cubic. A separate expanded sensitivity adds ArSSR, CycleINR, Pix2PixGAN, and SMamba-UNet only because all six aligned prediction sets exist; reconstruction fidelity was not used to select members.

All NIfTI identities, geometry, and generated indices are checked before scoring. The frozen configuration uses all non-target members, Eq. (2), raw inverse-normalized intensity, the 1% target mask, and 2,000 paired examination-bootstrap resamples. The cross-factor seed is 20260724; matched cases reuse bootstrap identities. Every legal factor-four committee subset is recomputed from voxel outputs. No external-test result changes the score, mask, or committee.

SwinIR flip-TTA is evaluated only where matching four-view maps exist. Pixel correlation is calculated per examination and aggregated by Fisher z; case ranking uses Spearman ρ with paired bootstrap intervals. Identical bootstrap indices are reused for paired comparisons. Mask variants that alter the foreground support are summarized by score-rank stability and endpoint-specific intervals rather than invalid paired differences against a changed MAE definition. Highest-MAE-decile AUROC/AUPRC is secondary because the two smaller cohorts contain only three positives per target. No profile or examination is removed after observing its score, and no sensitivity result replaces the frozen primary configuration on an external cohort.

3.2. Primary association and selective review

Table 2 shows moderate foreground localization (0.277–0.520) but strong examination ranking (0.890–0.966). Aggregation over generated foreground suppresses local noise and preserves between-case severity even when the map does not precisely identify each erroneous structure. nAURC is 0.024–0.092; at 50% retained coverage, referring

the highest-score half lowers retained MAE by 19.3–52.1% relative to random review. Confidence intervals are widest on fastMRI-brain and NingXia, so the continuous curves are more informative than a single threshold.

The complete curves in Fig. 2 avoid choosing a favorable test-set operating point. The largest absolute reduction occurs on IXI; smaller but consistent reductions occur on the two smaller cohorts. Because raw MAE is cohort-scale dependent, cross-cohort comparisons use nAURC and RRR rather than raw error.

3.3. Score design, committee membership, and preprocessing

Table 3. Score and committee ablations averaged over fixed target profiles (nAURC; lower is better).

Dataset	Score nAURC: med./mean/std.	Committee nAURC: neural/analytic/full
fastMRI-brain	0.073/0.070/0.074	0.017/0.118/0.073
IXI	0.032/0.043/0.049	0.033/0.059/0.032
NingXia	0.058/0.056/0.033	0.094/0.031/0.058

The median is a stable default but not a uniquely superior aggregator (Table 3). Across all 18 cross-factor profiles, its equal-profile ΔnAURC versus mean is -0.001 [-0.008, 0.008]; versus target-invariant all-output spread it is 0.011 [-0.007, 0.031] (negative favors median). Both intervals cross zero. The reliability signal is therefore not specific to the median, and target conditioning does not uniformly outperform generic examination difficulty. Committee composition is cohort dependent: neural peers are strongest on fastMRI-brain, analytical references on NingXia, and the full committee is near-best on IXI. Across legal factor-four subsets, nAURC spans 0.009–0.139; member identity matters more than size. The full committee is a frozen governance choice, not an empirical optimum.

Independent per-output normalization markedly degrades or reverses factor-four rankings, whereas input- and committee-derived masks largely preserve them. A common source-intensity representation is therefore part of the scoring contract; complete sensitivity results remain in the extended paper.

3.4. Shared-error boundary and qualitative behavior

The controlled stress test uses 20 fixed identity-sorted IXI examinations, all four targets, a foreground-constrained additive patch, five strengths, and a mask frozen before corruption. The target-specific arm perturbs only $\hat{y}_m$; the shared arm adds the identical patch to the target and every committee output.

At maximum strength, target MAE rises by 14.9% in both arms. Disagreement rises by 28.6% for target-specific corruption but changes by 0.0% under shared corruption, although mean committee MAE still rises by 11.8%. Figure 3 directly confirms Eq. (7); it is an identifiability test rather than a simulation of all clinical failures.

3.5. Cross-factor generalization and specificity

Reconstruction MAE increases monotonically from $\times 2$ to $\times 6$ in both cohorts, but review ranking does not track difficulty monotonically. Target-averaged fastMRI-brain ρ/nAURC improves from 0.888/0.098 to 0.937/0.052. NingXia peaks at $\times 4$ (0.926/0.058) and degrades at $\times 6$ (0.871/0.097), while RRR50 remains 20.9%; paired intervals do not establish degradation for every target. The complete plot and source data remain in the reproducibility package.

For direct target-specificity, we correlate each s_m with every target model's true error. The matched-minus-mismatched advantage is only -0.010–0.014 across six profiles, and every 95% interval includes zero. The score ranks difficult examinations well but does not preferentially identify which of the three neural targets failed; shared case difficulty dominates. This negative result narrows the interpretation of "target-to-committee" from empirical attribution to an explicitly conditioned scoring construction.

3.6. Single-model comparison and measured cost

On the same 24 fastMRI-brain $\times 4$ SwinIR examinations, four-view TTA yields $\rho = 0.960$ and nAURC 0.023 versus 0.931 and 0.061 for committee scoring. The paired nAURC difference is -0.038 [95% CI -0.119, 0.001], so the apparent gain is not conclusive. Stored-output auditing takes 1.18 s mean/2.87 s p95 per examination on CPU with 0.89 GiB peak RSS; online use additionally incurs every neural reconstruction.

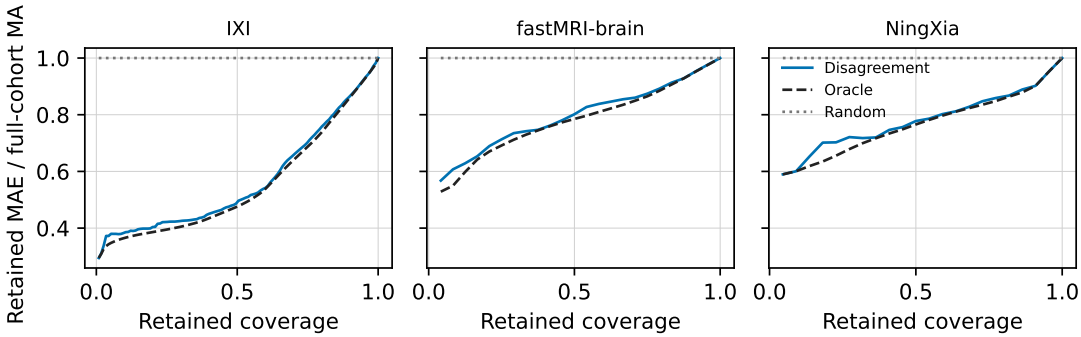


Fig. 2. Risk–coverage curves averaged over target interpolators after normalization by full-cohort MAE. The score remains between random and oracle review across cohorts; lower retained risk is better.

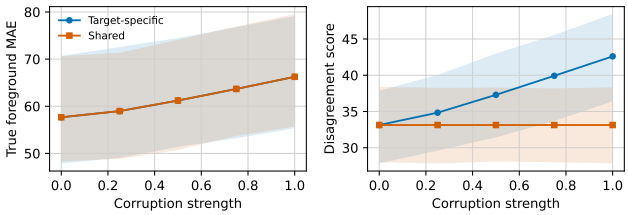


Fig. 3. Controlled shared-error test. Target-specific corruption raises both MAE and disagreement; identical corruption of all outputs raises MAE while cancelling from target-to-member differences. Bands are 95% bootstrap intervals.

4. EXTENDED ANALYSES

4.1. Registered candidate funnel and frozen validation

The universal-bank audit verifies 11 neural and two analytical methods across fastMRI-brain and NingXia at factors two, four, and six. I^{3}Net, HAT, and SwinIR are the universal targets because they also exist on IXI. Development uses IXI $\times 4$ and all fastMRI-brain factors; NingXia is not opened until four finalist definitions and their hash are serialized. Stage A evaluates 132 atomic score–mask–pooling candidates (74,052 case–target rows). Stage B combines five promoted definitions with eight committee policies, producing 40 score–committee candidates and 22,440 rows. Stage C uses 2,000 paired case-bootstrap resamples.

Table 4. Frozen finalists on development profiles. Worst is maximum profile nAURC.

Method	Mean nAURC	Worst	ρ	RRR50
Pre-specified	0.089	0.191	0.906	27.6%
Low-cost	0.087	0.159	0.906	27.8%
Dev-selected universal	0.040	0.099	0.947	28.1%
Shared-difficulty	0.036	0.093	0.952	28.1%

Table 5. One-shot NingXia OOF validation averaged over three targets.

Method	Factor	ρ	nAURC↓	RRR50↑
Pre-specified	$\times 2$	0.903	0.059	22.3%
Dev-selected universal	$\times 2$	0.797	0.130	18.4%
Low-cost	$\times 2$	0.912	0.059	21.5%
Shared-difficulty	$\times 2$	0.859	0.098	19.5%
Pre-specified	$\times 4$	0.927	0.058	22.2%
Dev-selected universal	$\times 4$	0.906	0.062	22.2%
Low-cost	$\times 4$	0.922	0.055	22.2%
Shared-difficulty	$\times 4$	0.899	0.075	22.2%
Pre-specified	$\times 6$	0.874	0.094	20.9%
Dev-selected universal	$\times 6$	0.812	0.154	19.2%
Low-cost	$\times 6$	0.846	0.127	19.9%
Shared-difficulty	$\times 6$	0.784	0.160	19.4%

The development-selected universal committee (I^{3}Net, SAINT, SMamba-UNet, and Pix2PixGAN, with target exclusion) cuts mean nAURC from 0.089 to 0.040. It fails to transfer: validation nAURC is 0.116 versus 0.070 for the pre-specified baseline, and RRR50 is 19.9% versus 21.8%. The paired nAURC difference is +0.049 [95% bootstrap interval $-0.038, 0.244$]. The low-cost analytical-anchor method is closer at 0.080 but is not demonstrably better. The shared/target decomposition chooses a fully shared score on development and also fails validation. No validation result is used to replace or revise a finalist.

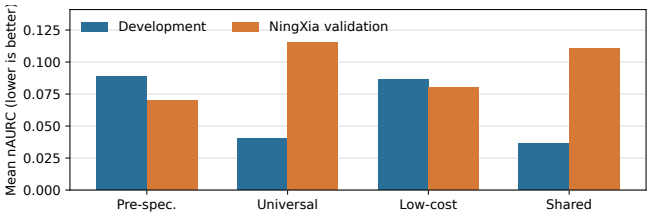


Fig. 4. Development-selected performance does not transfer to frozen NingXia validation. Lower nAURC is better.

Target-attribution matrices provide a second negative result. On NingXia, matched-minus-mismatched error correlation ranges from -0.008 to 0.014, usually with intervals containing zero. The construction is target-conditioned, but observed ranking mostly tracks shared examination difficulty. Cross-factor voxel consistency is not estimated because all 138 same-case comparisons have incompatible depth/generated support; defining a resampling rule after validation would violate the registered analysis.

4.2. Reconstruction fidelity and expanded diversity

Table 6 reports the original per-case reconstruction metrics from every model admitted to the expanded committee. All 42 dataset–factor–model rows were aggregated from the saved per-case CSV files; neither these metrics nor the resulting test errors were used to select the frozen primary committee. The added models span coordinate-based implicit reconstruction (ArSSR and CycleINR), adversarial convolutional synthesis (Pix2PixGAN), and a state-space/UNet hybrid (SMamba-UNet). Their availability allows a stricter diversity sensitivity than substituting one model at one factor.

Adding all four diverse networks changes mean nAURC from 0.098 to 0.096 at fastMRI-brain $\times 2$, from 0.073 to 0.088 at $\times 4$, and from 0.052 to 0.050 at $\times 6$. On NingXia, the corresponding changes are 0.065 to 0.110, 0.058 to 0.046, and 0.097 to 0.110. Thus broader architectural coverage is not monotonically beneficial: added members can sharpen consensus at one profile and dilute a useful analytical/neural contrast at another. This expanded analysis is intentionally labelled retrospective sensitivity and does not replace the five-member primary intersection.

The raw MAE scale is dataset specific, but within each cohort reconstruction difficulty rises with factor for every primary target. This trend does not force ranking quality to decline: reliability depends on whether difficult examinations create observable between-reconstructor separation, not only on mean reconstruction error. In particular, fastMRI-brain ranking improves as error grows, while NingXia $\times 6$ shows the clearest scale-dependent degradation.

4.3. Target specificity beyond shared case difficulty

For each dataset and factor, Fig. 5 reports

$$R_{m,b} = \mathrm{Spearman}(s_m, E_b), \tag{10}$$

$$A = \frac{1}{3}\sum_{m} R_{m,m} - \frac{1}{6}\sum_{m \neq b} R_{m,b}, \tag{11}$$

where A is the matched-target advantage. The matrix is fixed at 3×3 because I^{3}Net, HAT, and SwinIR are common neural targets at every profile. Paired case bootstrap resamples all nine cells with identical case indices.

Table 6. Reconstruction fidelity for the seven neural outputs available at every factor. Metrics use each run's saved per-case generated-slice evaluation. Values are descriptive and were not used for committee selection.

fastMRI-brain

Factor	Reconstructor	PSNR	SSIM	MAE	MSE	NCC
×2	ArSSR	23.11	0.761	43.6	8391.3	0.914
×2	CycleINR	23.10	0.759	43.7	8411.5	0.913
×2	HAT	24.65	0.811	37.3	6225.7	0.941
×2	I^3Net	27.77	0.867	25.3	3113.5	0.970
×2	Pix2PixGAN	26.07	0.830	31.2	4584.5	0.957
×2	SMamba-UNet	24.65	0.808	37.1	6361.1	0.940
×2	SwinIR	24.87	0.815	35.9	5990.4	0.943
×4	ArSSR	20.31	0.657	66.0	17745.9	0.830
×4	CycleINR	20.32	0.658	66.1	17743.2	0.830
×4	HAT	22.79	0.754	49.7	10454.5	0.908
×4	I^3Net	24.01	0.792	41.0	7755.7	0.931
×4	Pix2PixGAN	21.91	0.729	53.3	12485.2	0.891
×4	SMamba-UNet	22.73	0.762	49.4	10826.4	0.905
×4	SwinIR	22.36	0.745	52.0	11591.8	0.897
×6	ArSSR	19.42	0.616	74.9	21914.5	0.787
×6	CycleINR	19.39	0.616	74.9	22086.6	0.785
×6	HAT	21.82	0.731	55.8	13341.5	0.883
×6	I^3Net	22.83	0.762	47.1	10309.9	0.909
×6	Pix2PixGAN	21.65	0.725	54.8	13567.5	0.883
×6	SMamba-UNet	21.92	0.731	56.5	13166.7	0.885
×6	SwinIR	21.84	0.731	56.4	13267.8	0.884

NingXia

Factor	Reconstructor	PSNR	SSIM	MAE	MSE	NCC
×2	ArSSR	25.16	0.747	78.8	29980.4	0.911
×2	CycleINR	25.16	0.747	78.8	29980.4	0.911
×2	HAT	25.85	0.777	71.8	25723.3	0.924
×2	I^3Net	25.75	0.761	73.0	26688.8	0.923
×2	Pix2PixGAN	25.30	0.760	75.9	27897.9	0.915
×2	SMamba-UNet	25.45	0.764	76.7	28039.5	0.917
×2	SwinIR	25.58	0.768	75.3	27450.9	0.919
×4	ArSSR	23.21	0.714	98.3	41837.5	0.870
×4	CycleINR	22.94	0.698	103.3	44475.7	0.861
×4	HAT	23.99	0.755	87.8	35193.7	0.892
×4	I^3Net	23.48	0.724	96.1	39453.8	0.881
×4	Pix2PixGAN	22.44	0.716	108.0	51103.2	0.852
×4	SMamba-UNet	23.96	0.754	88.3	35445.0	0.891
×4	SwinIR	24.01	0.756	87.3	34871.0	0.892
×6	ArSSR	22.39	0.704	110.8	51461.2	0.836
×6	CycleINR	21.93	0.676	119.1	57033.3	0.818
×6	HAT	23.07	0.740	100.5	44299.8	0.862
×6	I^3Net	22.40	0.695	115.0	51157.7	0.843
×6	Pix2PixGAN	21.25	0.673	130.7	65474.6	0.801
×6	SMamba-UNet	23.01	0.737	102.8	45207.3	0.860
×6	SwinIR	23.04	0.740	100.6	44639.5	0.861

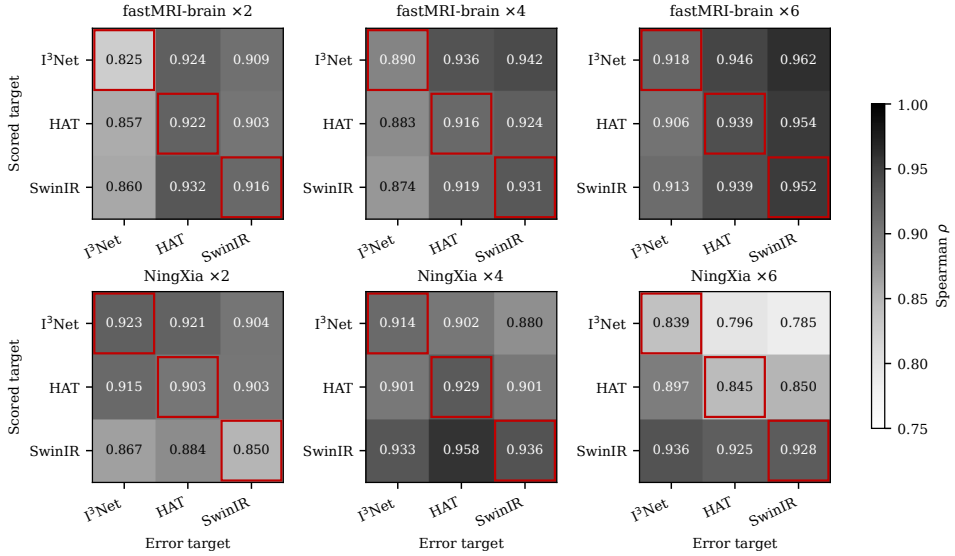


Fig. 5. Direct target-specificity test. Rows identify the score's named target and columns the model whose true error is correlated with that score. Red boxes mark matched cells. Similar within-row values show that shared examination difficulty dominates target identity.

Table 7. Matched-target advantage from Eq. (11). Every interval includes zero.

Cohort	Factor	Matched advantage [95% CI]	Row wins
fastMRI-brain	×2	-0.010 [-0.036, 0.008]	2/3
fastMRI-brain	×4	-0.001 [-0.013, 0.008]	2/3
fastMRI-brain	×6	-0.000 [-0.017, 0.018]	2/3
NingXia	×2	-0.007 [-0.039, 0.007]	1/3
NingXia	×4	0.014 [-0.001, 0.051]	2/3
NingXia	×6	0.006 [-0.019, 0.033]	1/3

The six advantages range from -0.010–0.014, with only one or two of three diagonal cells exceeding their row's mismatched mean. The target-invariant all-output spread has exactly zero advantage by construction. Raw within-examination top-1 attribution is also unstable: it is 0% on all fastMRI factors, 36.4% on NingXia ×2, and 95.5% on NingXia ×4/×6. Score magnitudes therefore should not be compared across target identities without calibration. These results do not invalidate examination-level selective review; they invalidate the stronger interpretation that the score identifies which reconstructor is uniquely responsible for a difficult case.

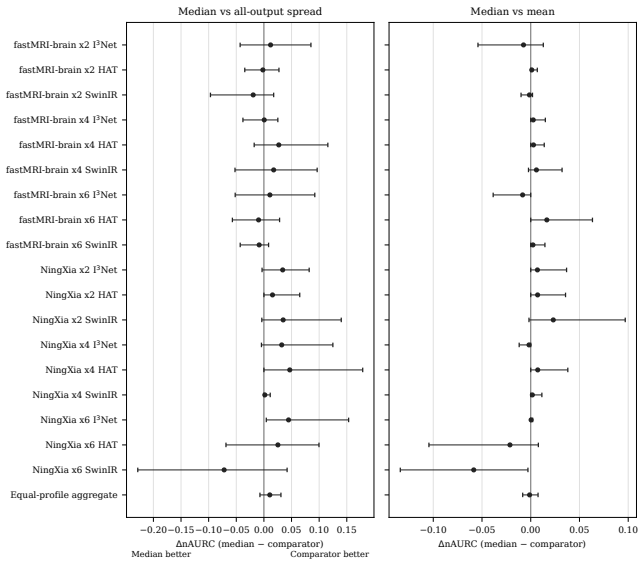


Fig. 6. Paired nAURC differences for every fixed cross-factor profile and the equal-profile aggregate. Intervals are 95% paired case-bootstrap intervals. No profile was removed after observing its comparison.

Table 8. Equal-profile paired comparisons; negative ΔnAURC favors median.

Comparator	ΔnAURC [95% CI]	Profiles
Mean target-to-member	-0.001 [-0.008, 0.008]	18
All-output spread	0.011 [-0.007, 0.031]	18
Excluded-committee spread	-0.003 [-0.026, 0.020]	18

4.4. Paired comparator analysis

Figure 6 uses the same bootstrap case identities for both members of each comparison. For method A equal to target median and method B the comparator, $\Delta\mathrm{nAURC} = \mathrm{nAURC}_A - \mathrm{nAURC}_B$; negative values favor the median. Equal-profile aggregates average the 18 fixed dataset–factor–target endpoints rather than pooling examinations across cohorts.

Median versus mean is effectively unresolved at the aggregate level (-0.001 [-0.008, 0.008]). Median is also not superior to all-output spread (0.011 [-0.007, 0.031]), and its difference from target-excluded spread is -0.003 [95% CI -0.026, 0.020]. Full versus neural-only and analytical-only comparisons likewise vary by cohort and factor. We therefore retain the median as a pre-specified robust default, not as an empirically dominant algorithmic component. The contribution is the observable-output audit protocol, strict target exclusion, generated-support contract, selective-review evaluation, and explicit identification of what the score cannot establish.

4.5. Sensitivity, failure examples, and deployment detail

Independent p1–p99, median/IQR, or min–max normalization can destroy cross-model amplitude information and substantially worsen ranking. In contrast, committee-union, committee-intersection, committee-median, and observed-input masks mostly preserve examination ordering; high target thresholds are least stable on fastMRI-brain. The complete factor-four summary is retained in Table 9.

Table 9. Factor-four sensitivity and secondary checks. Normalization columns report deterioration from raw intensity; mask entries are score-rank stability.

Analysis	Result	Interpretation
Normalization	ΔnAURC $+0.248$–$+1.359$	independent scaling degrades ranking
Mask stability	min. $\rho_s = 0.606$; most ≥ 0.973	robust except high fastMRI threshold
Legal subsets	nAURC 0.009–0.139	member identity matters more than size
SwinIR TTA	0.023 vs. 0.061 nAURC	local gain inconclusive by paired CI
CPU audit	1.18 s mean / 2.87 s p95	stored-output post-processing only

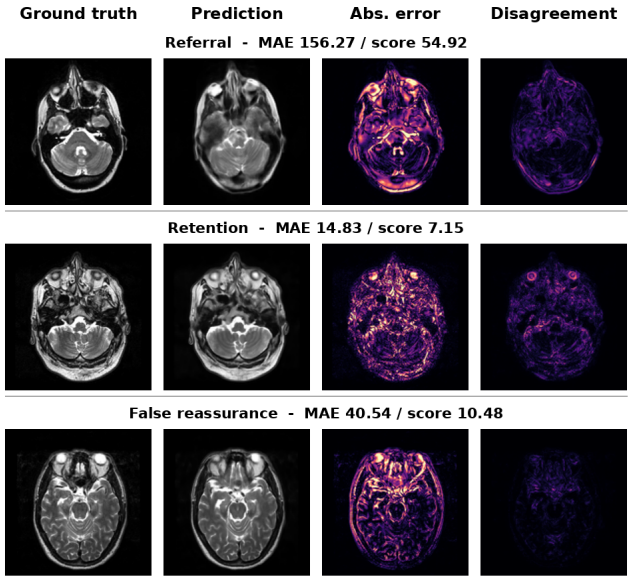


Fig. 7. Automatically selected IXI/SwinIR examples: useful referral, stable retention, and false reassurance. Error is not interpreted as diagnostic harm.

Automatically selected failure examples complement the controlled common-mode experiment: a high-error examination can receive a low score when peers share over-smoothing, while a correct but unconventional detail can trigger unnecessary review. Highest-MAE-decile AUROC/AUPRC remains secondary because fastMRI-brain and NingXia contain only three positive examinations per target. Online deployment must additionally benchmark every neural member on matched hardware; the measured 1.18-s post-hoc time applies only when outputs already exist.

5. DISCUSSION AND CONCLUSION

The evidence supports examination triage more strongly than precise voxel localization. A high s_m prioritizes inspection while preserving the original reconstruction and committee identities; a low score indicates consensus only within the available committee. Coverage should follow review capacity rather than a retrospectively tuned error threshold, and a deployment threshold requires validation after changes in scanner, reconstruction software, intensity mapping, or committee membership.

The protocol is most attractive for retrospective audit, benchmark comparison, or systems that already cache several deterministic outputs. If only one model can be run, TTA, posterior sampling, acquisition residuals, or a validated error predictor may be more economical. A larger committee is not automatically better: member identity matters more than count, and the available no-reference

diversity proxy cannot reliably choose the best compact subset. In practice, a review system should therefore freeze the committee specification, log member identities with every score, and monitor score distributions after software or acquisition changes rather than silently swapping members.

The score is also complementary to model-internal uncertainty. TTA probes sensitivity of one model to transformations, whereas cross-model disagreement probes separation from alternative reconstructions. Their failure modes need not coincide: one model may be unstable while peers agree, or all models may be confidently wrong together. A future deployment study should measure whether presenting both signals improves reviewer prioritization, review time, or downstream robustness without encouraging automation bias.

The principal limitations are structural. Shared errors are invisible, correlated members reduce diversity, a target mask can omit suppressed tissue, and unusual correct detail can be over-referred. Evidence remains restricted to T2 brain MRI; fastMRI-brain and NingXia are small, $\times 6$ NingXia degrades, and the matched analysis cannot separate target identity from shared difficulty. There is no reader study or downstream task, and MAE is not diagnostic harm. Target-to-committee disagreement is therefore a reproducible ranking layer for selective review of stored outputs, not calibrated uncertainty, model attribution, or a clinical guarantee.

6. REFERENCES

[1] Haofei Song, Xintian Mao, Jing Yu, Qingli Li, and Yan Wang, "I^3Net: Inter-intra-slice interpolation network for medical slice synthesis," *IEEE Transactions on Medical Imaging*, 2024.

[2] Jingyun Liang, Jiezhang Cao, Guolei Sun, Kai Zhang, Luc Van Gool, and Radu Timofte, "SwinIR: Image restoration using swin transformer," in *Proceedings of the IEEE/CVF International Conference on Computer Vision Workshops*, 2021.

[3] Xiangyu Chen, Xintao Wang, Jiantao Zhou, Yu Qiao, and Chao Dong, "Activating more pixels in image super-resolution transformer," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2023.

[4] Yarin Gal and Zoubin Ghahramani, "Dropout as a bayesian approximation: Representing model uncertainty in deep learning," in *Proceedings of the International Conference on Machine Learning*, 2016.

[5] Balaji Lakshminarayanan, Alexander Pritzel, and Charles Blundell, "Simple and scalable predictive uncertainty estimation using deep ensembles," in *Advances in Neural Information Processing Systems*, 2017.

[6] Alex Kendall and Yarin Gal, "What uncertainties do we need in bayesian deep learning for computer vision?," in *Advances in Neural Information Processing Systems*, 2017.

[7] Guotai Wang, Wenqi Li, Michael Aertsen, Jan Deprest, Sebastien Ourselin, and Tom Vercauteren, "Aleatoric uncertainty estimation with test-time augmentation for medical image segmentation with convolutional neural networks," *arXiv preprint arXiv:1807.07356*, 2018.

[8] Vineet Edupuganti, Morteza Mardani, Shreyas Vasanawala, and John M. Pauly, "Uncertainty quantification in deep MRI reconstruction," *IEEE Transactions on Medical Imaging*, vol. 40, no. 1, pp. 239–250, 2021.

[9] Thomas Kuestner, Kerstin Hammernik, Daniel Rueckert, Tobias Hepp, and Sergios Gatidis, "Predictive uncertainty in deep learning-based MR image reconstruction using deep ensembles: Evaluation on the fastMRI data set," *Magnetic Resonance in Medicine*, vol. 92, no. 1, pp. 289–302, 2024.

[10] Richard Shaw, Carole H. Sudre, Sebastien Ourselin, and M. Jorge Cardoso, "A heteroscedastic uncertainty model for decoupling sources of MRI image quality," in *Proceedings of the Third Conference on Medical Imaging with Deep Learning*. 2020, vol. 121 of *Proceedings of Machine Learning Research*, pp. 733–742, PMLR.

[11] Shi Hu, Nicola Pezzotti, and Max Welling, "Learning to predict error for MRI reconstruction," in *Medical Image Computing and Computer Assisted Intervention – MICCAI 2021*. 2021, vol. 12903 of *Lecture Notes in Computer Science*, pp. 604–613, Springer.

[12] Jiaxin Zhang, Sirui Bi, and Victor Fung, "On the quantification of image reconstruction uncertainty without training data," in *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, 2024, pp. 2072–2081.

[13] Igor Stepien, Rafal Obuchowicz, Katarzyna Pieszko, Paul Paciorek, and Marcin Kociolek, "A brief survey on no-reference image quality assessment methods for magnetic resonance images," *Journal of Imaging*, vol. 8, no. 6, pp. 160, 2022.

[14] Ran El-Yaniv and Yair Wiener, "On the foundations of noise-free selective classification," *Journal of Machine Learning Research*, vol. 11, no. 53, pp. 1605–1641, 2010.

[15] Yonatan Geifman and Ran El-Yaniv, "Selectivenet: A deep neural network with an integrated reject option," in *Proceedings of the International Conference on Machine Learning*, 2019.

[16] Yukun Ding, Jinglan Liu, Jinjun Xiong, and Yiyu Shi, "Revisiting the evaluation of uncertainty estimation and its application to explore model complexity–uncertainty trade-off," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops*, 2020.

[17] Benjamin Kompa, Jasper Snoek, and Andrew L. Beam, "Second opinion needed: Communicating uncertainty in medical machine learning," *npj Digital Medicine*, vol. 4, no. 1, pp. 4, 2021.

[18] Imperial College London, "IXI dataset," 2010, Accessed as a public neuroimaging dataset.

[19] Jure Zbontar, Florian Knoll, Anuroop Sriram, Tullie Murrell, Zhengnan Huang, Matthew J. Muckley, Aaron Defazio, Ruben Stern, Patricia Johnson, Mary Bruno, Marc Parente, Krzysztof J. Geras, Joe Katsnelson, Hersh Chandarana, Zizhao Zhang, Michal Drozdzal, Adriana Romero, Michael Rabbat, Pascal Vincent, Nafissa Yakubova, James Pinkerton, Duo Wang, Erich Owens, C. Lawrence Zitnick, Michael P. Recht, Daniel K. Sodickson, and Yvonne W. Lui, "fastMRI: An open dataset and benchmarks for accelerated MRI," *arXiv preprint arXiv:1811.08839*, 2018.

[20] Robert G. Keys, "Cubic convolution interpolation for digital image processing," *IEEE Transactions on Acoustics, Speech, and Signal Processing*, vol. 29, no. 6, pp. 1153–1160, 1981.

[21] Philippe Thévenaz, Thierry Blu, and Michael Unser, "Interpolation revisited," *IEEE Transactions on Medical Imaging*, vol. 19, no. 7, pp. 739–758, 2000.

[22] Michael Unser, "Splines: A perfect fit for signal and image processing," *IEEE Signal Processing Magazine*, vol. 16, no. 6, pp. 22–38, 1999.